

## Problem - 1: Perform a classification task with knn from scratch.

### 1. Load the Dataset:

- Read the dataset into a pandas DataFrame.
- Display the first few rows and perform exploratory data analysis (EDA) to understand the dataset (e.g., check data types, missing values, summary statistics).

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import time

df = pd.read_csv("C:\everything\herald\AI\week4\diabetes_.csv")
print("\nFirst 5 rows:")
print(df.head())

print("\nData types: ")
print(df.dtypes)

print("\nMissing values: ")
print(df.isna().sum())

print("\nSummary statistics: ")
print(df.describe(include='all'))
```

```

First 5 rows:
   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI \
0            6      148           72          35        0  33.6
1            1       85           66          29        0  26.6
2            8      183           64           0        0  23.3
3            1       89           66          23        94 28.1
4            0      137           40          35       168 43.1

   DiabetesPedigreeFunction  Age  Outcome
0                  0.627    50       1
1                  0.351    31       0
2                  0.672    32       1
3                  0.167    21       0
4                 2.288    33       1

Data types:
Pregnancies           int64
Glucose               int64
BloodPressure         int64
SkinThickness         int64
Insulin              int64
BMI                  float64
DiabetesPedigreeFunction float64
Age                  int64
Outcome              int64
dtype: object

Missing values:
Pregnancies      0
Glucose          0
BloodPressure    0
SkinThickness    0
Insulin          0
BMI              0
DiabetesPedigreeFunction 0
Age              0
Outcome          0
dtype: int64

Summary statistics:
   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin \
count    768.000000  768.000000  768.000000  768.000000  768.000000
mean     3.845052 120.894531  69.105459  20.536458 79.79479
std      3.369578 31.972618 19.355807 15.952218 115.244002
min      0.000000  0.000000  0.000000  0.000000  0.000000
25%     1.000000  99.000000  62.000000  0.000000  0.000000
50%     3.000000 117.000000  72.000000 23.000000 30.500000
75%     6.000000 140.250000  80.000000 32.000000 127.250000
max     17.000000 199.000000 122.000000 99.000000 846.000000

   BMI  DiabetesPedigreeFunction  Age  Outcome
count  768.000000             768.000000  768.000000  768.000000
mean   31.992578             0.471876 33.240885  0.348958
std    7.884160              0.331329 11.760232  0.476951
min    0.000000              0.078000 21.000000  0.000000
25%   27.300000              0.243750 24.000000  0.000000
50%   32.000000              0.372500 29.000000  0.000000
75%   36.500000              0.626250 41.000000  1.000000
max   67.100000              2.428000 81.000000  1.000000

```

## 2. Handle Missing Data:

- Handle any missing values appropriately, either by dropping or imputing them based on the data.

```
df_handle = df.copy()
for col in df_handle.columns:
    if df_handle[col].isna().any():
        median_value = df_handle[col].median() if np.issubdtype(df_handle[col].dtype, np.number) else df_handle[col].mode()[0]
        df_handle[col] = df_handle[col].fillna(median_value)

print("Any remaining NaNs?", df_handle.isna().any().any())
Any remaining NaNs? False
```

## 3. Feature Engineering:

- Separate the feature matrix (X) and target variable (y).
- Perform a train - test split from scratch using a 70% – 30% ratio.

```
X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values

np.random.seed(42)
indices = np.random.permutation(len(X))

train_size = int(0.7 * len(X))
train_idx = indices[:train_size]
test_idx = indices[train_size:]

X_train = X[train_idx]
X_test = X[test_idx]
y_train = y[train_idx]
y_test = y[test_idx]
print("Train size:", X_train.shape[0], "Test size:", X_test.shape[0])
```

Train size: 537 Test size: 231

#### 4. Implement KNN:

- Build the KNN algorithm from scratch (no libraries like sickit-learn for KNN).
- Compute distances using Euclidean distance.
- Write functions for:
  - Predicting the class for a single query.
  - Predicting classes for all test samples.
- Evaluate the performance using accuracy.

```
def euclidean_distance(a: np.ndarray, b: np.ndarray):
    return np.sqrt(np.sum(a - b) ** 2)

def knn_predict_one(x_query: np.ndarray, X_train: np.ndarray, y_train: np.ndarray, k: int = 5):
    dists = np.sqrt(np.sum((X_train - x_query) ** 2, axis=1))
    nn_idx = np.argpartition(dists, k)[:k]
    nn_labels = y_train[nn_idx]
    values, counts = np.unique(nn_labels, return_counts=True)
    return values[np.argmax(counts)]

def knn_predict(X_query: np.ndarray, X_train: np.ndarray, y_train: np.ndarray, k: int = 5):
    preds = np.zeros(X_query.shape[0], dtype=y_train.dtype)
    for i in range(X_query.shape[0]):
        preds[i] = knn_predict_one(X_query[i], X_train, y_train, k)
    return preds

def accuracy_score(y_true: np.ndarray, y_pred: np.ndarray):
    return (y_true == y_pred).mean()

k_baseline = 5
t0 = time.time()
y_pred_baseline = knn_predict(X_test, X_train, y_train, k=k_baseline)
t1 = time.time()

acc_baseline = accuracy_score(y_test, y_pred_baseline)
print(f"Baseline (unscaled) KNN | k={k_baseline}: Accuracy={(acc_baseline:.4f)}, Time={(t1 - t0:.4f)s}")

Baseline (unscaled) KNN | k=5: Accuracy=0.7143, Time=0.0190s
```

## Problem - 2 - Experimentation:

### 1. Repeat the Classification Task:

- Scale the Feature matrix X.
- Use the scaled data for training and testing the kNN Classifier.
- Record the results.

```
#problem 2
X_mean = X_train.mean(axis=0)
X_std = X_train.std(axis=0)
X_std[X_std == 0] = 1.0

X_train_scaled = (X_train - X_mean) / X_std
X_test_scaled = (X_test - X_mean) / X_std

k_scaled = 3
t0 = time.time()
y_pred_scaled = knn_predict(X_test_scaled, X_train_scaled, y_train, k=k_scaled)
t1 = time.time()

acc_scaled = accuracy_score(y_test, y_pred_scaled)
print(f"Scaled KNN | k={k_scaled}: Accuracy={acc_scaled:.4f}, Time={(t1 - t0):.4f}s")

Scaled KNN | k=3: Accuracy=0.6970, Time=0.0158s
```

### 2. Comparative Analysis: Compare the Results -

- Compare the accuracy and performance of the kNN model on the original dataset from problem 1 versus the scaled dataset.
- Discuss:
  - How scaling impacted the KNN performance.
  - The reason for any observed changes in accuracy.

```
print("Comparison at k=3")
print(f"- Unscaled accuracy: {acc_baseline:.4f}")
print(f"- Scaled accuracy: {acc_scaled:.4f}")
print("Observation: Scaling changes feature magnitudes, which can impact neighbor selection and thus accuracy.")

Comparison at k=3
- Unscaled accuracy: 0.7143
- Scaled accuracy: 0.6970
Observation: Scaling changes feature magnitudes, which can impact neighbor selection and thus accuracy.
```

## Problem - 3 - Experimentation with k:

### 1. Vary the number of neighbors - k:

- Run the KNN model on both the original and scaled datasets for a range of: k= 1, 2, 3, . . . 15

- For each k, record:
  - Accuracy.
  - Time taken to make predictions.

```
#problem 3
ks = list(range(1, 16))
acc_unscaled_list, time_unscaled_list = [], []
acc_scaled_list, time_scaled_list = [], []

for k in ks:
    t0 = time.time()
    y_pred_u = knn_predict(X_test, X_train, y_train, k=k)
    t1 = time.time()
    acc_unscaled_list.append(accuracy_score(y_test, y_pred_u))
    time_unscaled_list.append(t1 - t0)

    t0 = time.time()
    y_pred_s = knn_predict(X_test_scaled, X_train_scaled, y_train, k=k)
    t1 = time.time()
    acc_scaled_list.append(accuracy_score(y_test, y_pred_s))
    time_scaled_list.append(t1 - t0)

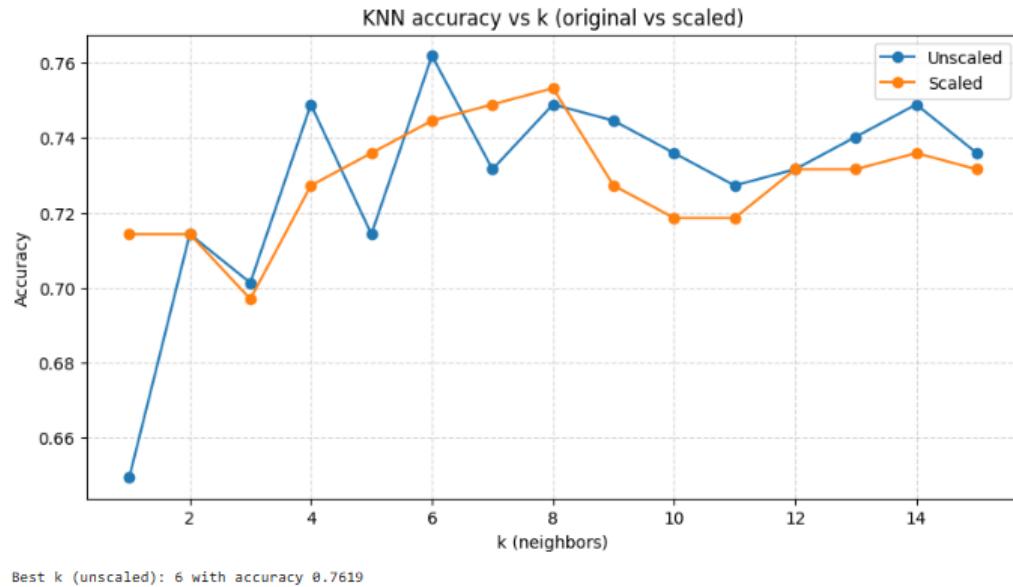
best_k_unscaled = ks[int(np.argmax(acc_unscaled_list))]
best_k_scaled = ks[int(np.argmax(acc_scaled_list))]
```

## 2. Visualize the Results:

- Plot the following graphs:
  - k vs. Accuracy for original and scaled datasets.
  - k vs. Time Taken for original and scaled datasets

```
plt.figure(figsize=(10, 5))
plt.plot(ks, acc_unscaled_list, marker='o', label='Unscaled')
plt.plot(ks, acc_scaled_list, marker='o', label='Scaled')
plt.xlabel('k (neighbors)')
plt.ylabel('Accuracy')
plt.title('KNN accuracy vs k (original vs scaled)')
plt.grid(True, linestyle='--', alpha=0.5)
plt.legend()
plt.show()

print(f"Best k (unscaled): {best_k_unscaled} with accuracy {max(acc_unscaled_list):.4f}")
print(f"Best k (scaled): {best_k_scaled} with accuracy {max(acc_scaled_list):.4f}")
```



### .3. Analyze and Discuss:

- Discuss how the choice of k affects the accuracy and computational cost.
- Identify the optimal k based on your analysis.

```
plt.figure(figsize=(10, 5))
plt.plot(ks, time_unscaled_list, marker='o', label='Unscaled')
plt.plot(ks, time_scaled_list, marker='o', label='Scaled')
plt.xlabel('k (neighbors)')
plt.ylabel('Prediction time (s)')
plt.title('KNN prediction time vs k (original vs scaled)')
plt.grid(True, linestyle='--', alpha=0.5)
plt.legend()
plt.show()
```

