

Part 1: Regression Task (California Housing)

Task 1: Load and Split Dataset

```
[1]:  
import pandas as pd  
import numpy as np  
from sklearn.model_selection import train_test_split  
  
# Setting the columns  
cols = ["longitude", "latitude", "housingMedianAge", "totalRooms", "totalBedrooms", "population", "households", "medianIncome", "medianHouseValue"]  
  
# Loading the dataset by downloading from "https://s3-eu-west-1.amazonaws.com/pfigshare-u-files/5976036/cal_housing.tgz"  
df = pd.read_csv("cal_housing.data", header=None, names=cols)  
  
# Splitting the data into features and label  
X = df.drop("medianHouseValue", axis=1)  
y = df["medianHouseValue"]  
  
# Split into training (80%) and test (20%)  
X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.2, random_state=42  
)  
  
print(f"\nTraining set: {X_train.shape}")  
print(f"Test set: {X_test.shape}")  
print(f"Features: {X_train.shape[1]}")  
  
Training set: (16512, 8)  
Test set: (4128, 8)  
Features: 8
```

Task 2, Step 1: Baseline Model (No Regularization)

```
[2]:  
from sklearn.linear_model import LinearRegression  
from sklearn.metrics import mean_squared_error  
  
# Build Linear Regression model  
model = LinearRegression()  
model.fit(X_train, y_train)  
  
# Observe coefficients  
print("Coefficients:")  
for i in range(min(5, len(model.coef_))): # Show first 5  
    print(f" Feature {i}: {model.coef_[i]:.6f}")  
print(f" ... and {len(model.coef_) - 5} more")  
print(f"Intercept: {model.intercept_:.6f}")  
  
# Compute MSE  
y_train_pred = model.predict(X_train)  
y_test_pred = model.predict(X_test)  
  
mse_train = mean_squared_error(y_train, y_train_pred)  
mse_test = mean_squared_error(y_test, y_test_pred)  
  
print(f"\nTraining MSE: {mse_train:.6f}")  
print(f"Test MSE: {mse_test:.6f}")
```

Coefficients:
Feature 0: -42632.391717
Feature 1: -42458.071863
Feature 2: 1182.809649
Feature 3: -8.187977
Feature 4: 116.260128
... and 3 more
Intercept: -3578224.234818

Training MSE: 4811134397.884196
Test MSE: 4918556441.477828

Task 2, Step 2: Hyperparameter Tuning

```
[3]:  
from sklearn.linear_model import Ridge, Lasso  
from sklearn.model_selection import GridSearchCV  
  
# Define alphas  
alphas = [0.001, 0.01, 0.1, 1, 10, 100, 1000]  
  
# Ridge Regression tuning  
ridge = Ridge()  
ridge_grid = GridSearchCV(ridge, {'alpha': alphas}, cv=5, scoring='neg_mean_squared_error')  
ridge_grid.fit(X_train, y_train)  
  
print(f"Ridge - Best alpha: {ridge_grid.best_params_['alpha']}")  
  
# Lasso Regression tuning  
lasso = Lasso(max_iter=10000)  
lasso_grid = GridSearchCV(lasso, {'alpha': alphas}, cv=5, scoring='neg_mean_squared_error')  
lasso_grid.fit(X_train, y_train)  
  
print(f"Lasso - Best alpha: {lasso_grid.best_params_['alpha']}")  
  
# Test set evaluation  
ridge_pred = ridge_grid.predict(X_test)  
lasso_pred = lasso_grid.predict(X_test)
```

```

print(f"\nLasso - Best alpha: {lasso_grid.best_params_['alpha']}")

# Test set evaluation
ridge_pred = ridge_grid.predict(X_test)
lasso_pred = lasso_grid.predict(X_test)

ridge_mse = mean_squared_error(y_test, ridge_pred)
lasso_mse = mean_squared_error(y_test, lasso_pred)

print(f"\nTest MSE - Ridge: {ridge_mse:.6f}")
print(f"Test MSE - Lasso: {lasso_mse:.6f}")

Ridge - Best alpha: 10
Lasso - Best alpha: 10

Test MSE - Ridge: 4918567284.465969
Test MSE - Lasso: 4918555581.562371

```

Task 2, Step 3: Regularization Experiments (L1 vs L2)

```

[4]: # Train models with best parameters
ridge_best = Ridge(alpha=ridge_grid.best_params_['alpha'])
lasso_best = Lasso(alpha=lasso_grid.best_params_['alpha'], max_iter=10000)

ridge_best.fit(X_train, y_train)
lasso_best.fit(X_train, y_train)

# Compare coefficients
print("\nCoefficient Comparison (first 8 features):")
print("Feature\tBaseline\t\tRidge\t\t\tLasso")
for i in range(8):
    print(f"{i}\t{(model.coef_[i]):.6f}\t{ridge_best.coef_[i]:.6f}\t{lasso_best.coef_[i]:.6f}")

# Count zero coefficients
zero_lasso = sum(lasso_best.coef_ == 0)
print(f"\nZero coefficients in Lasso: {zero_lasso}/{len(lasso_best.coef_)}")

# Performance comparison
ridge_train_mse = mean_squared_error(y_train, ridge_best.predict(X_train))
ridge_test_mse = mean_squared_error(y_test, ridge_best.predict(X_test))

lasso_train_mse = mean_squared_error(y_train, lasso_best.predict(X_train))
lasso_test_mse = mean_squared_error(y_test, lasso_best.predict(X_test))

print(f"\nPerformance Comparison:")
print(f"{'Model':<10} {'Train MSE':<15} {'Test MSE':<15}")
print("-" * 40)
print(f"{'Baseline':<10} {mse_train:<15.6f} {mse_test:<15.6f}")
print(f"{'Ridge':<10} {(ridge_train_mse:<15.6f)} {(ridge_test_mse:<15.6f)}")
print(f"{'Lasso':<10} {(lasso_train_mse:<15.6f)} {(lasso_test_mse:<15.6f)}")

print("\nDiscussion:")
print("1. L1 produces sparse coefficients (feature selection)")
print("2. L2 shrinks coefficients without zeroing them")
print("3. Regularization reduces variance, prevents overfitting")
print("4. Excessive regularization increases bias")

```

Coefficient Comparison (first 8 features):				
	Feature	Baseline	Ridge	Lasso
0		-42632.391717	-42535.627082	-42595.282794
1		-42450.071863	-42359.666504	-42415.402048
2		1182.889649	1184.351988	1183.328980
3		-8.187977	-8.196936	-8.191118

```

# Performance comparison
ridge_train_mse = mean_squared_error(y_train, ridge_best.predict(X_train))
ridge_test_mse = mean_squared_error(y_test, ridge_best.predict(X_test))

lasso_train_mse = mean_squared_error(y_train, lasso_best.predict(X_train))
lasso_test_mse = mean_squared_error(y_test, lasso_best.predict(X_test))

print("\nPerformance Comparison:")
print(f"\tModel: {10} \tTrain MSE: {15} \tTest MSE: {15}")
print("-" * 48)
print(f"\tBaseline: {10} \t{mse_train:{15.6f}} \t{mse_test:{15.6f}}")
print(f"\tRidge: {10} \t{ridge_train_mse:{15.6f}} \t{ridge_test_mse:{15.6f}}")
print(f"\tLasso: {10} \t{lasso_train_mse:{15.6f}} \t{lasso_test_mse:{15.6f}}")

print("\nDiscussion:")
print("1. L1 produces sparse coefficients (feature selection)")
print("2. L2 shrinks coefficients without zeroing them")
print("3. Regularization reduces variance, prevents overfitting")
print("4. Excessive regularization increases bias")

```

```

Coefficient Comparison (first 8 features):
Feature Baseline          Ridge          Lasso
8      -42632.391717    -42535.627082    -42595.282794
1      -42458.071863    -42359.666584    -42415.482048
2      1182.809649     1184.351988     1183.328988
3      -8.187977     -8.196936     -8.191118
4      116.260128     116.124492     116.204998
5      -38.492213     -38.496151     -38.493883
6      46.342572      46.569026      46.438671
7      40538.404387    40543.565513    40540.088671

Zero coefficients in Lasso: 0/8

Performance Comparison:
Model      Train MSE      Test MSE
-----
Baseline   4811134397.884196 4918556441.477828
Ridge      4811139082.000712 4918567284.465969
Lasso      4811135093.259236 4918555581.562371

Discussion:
1. L1 produces sparse coefficients (feature selection)
2. L2 shrinks coefficients without zeroing them
3. Regularization reduces variance, prevents overfitting
4. Excessive regularization increases bias

```

PART 2: CLASSIFICATION TASK (Breast Cancer)

Task 1: Load and Split Dataset

```

[5]:
from sklearn.datasets import load_breast_cancer

X, y = load_breast_cancer(return_X_y=True)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

print(f"Training set size: {X_train.shape}")
print(f"Test set size: {X_test.shape}")

Training set size: (455, 30)
Test set size: (114, 30)

```

```
Test set size: (114, 30)
```

Task 2, Step 1: Baseline Model (No Regularization)

```
[8]: # Baseline Logistic Regression
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

log_reg = LogisticRegression(max_iter=10000)
log_reg.fit(X_train, y_train)

# Predictions
y_train_pred = log_reg.predict(X_train)
y_test_pred = log_reg.predict(X_test)

# Accuracy
print("Baseline Logistic Regression")
print("Train Accuracy:", accuracy_score(y_train, y_train_pred))
print("Test Accuracy:", accuracy_score(y_test, y_test_pred))

# Coefficients
print("Coefficients:", log_reg.coef_)

Baseline Logistic Regression
Train Accuracy: 0.0626373626373627
Test Accuracy: 0.956140358877193
Coefficients: [[ 0.98744559  0.22584768 -0.36798308  0.02622024 -0.15343354 -0.23442582
 -0.52264211 -0.27397911 -0.22365307 -0.03748764 -0.09513651  1.39324606
 -0.16864384 -0.08886037 -0.02191314  0.04244395 -0.048801058 -0.03160105
 -0.03424751  0.01082329  0.8974889 -0.51492842 -0.01604828 -0.01662147
 -0.30310642 -0.77708659 -1.42290989 -0.49745618 -0.73363954 -0.10287671]]
```

Task 2, Step 2: Hyperparameter Tuning

```
[9]: param_grid = {'C': [0.001, 0.01, 0.1, 1, 10, 100], 'solver': ['liblinear']}
log_cv = GridSearchCV(log_reg, param_grid, scoring = 'accuracy', cv = 5)
log_cv.fit(X_train, y_train)

print("Best Parameters:", log_cv.best_params_)

best_log = log_cv.best_estimator_

y_train_pred = best_log.predict(X_train)
y_test_pred = best_log.predict(X_test)

print("Tuned Logistic Regression")
print("Train Accuracy:", accuracy_score(y_train, y_train_pred))
print("Test Accuracy:", accuracy_score(y_test, y_test_pred))

# L1 Logistic Regression
log_l1 = LogisticRegression(C=log_cv.best_params_['C'], l1_ratio = 1.0, solver='liblinear', max_iter=10000)
log_l1.fit(X_train, y_train)

# L2 Logistic Regression
log_l2 = LogisticRegression(C=log_cv.best_params_['C'], l1_ratio = 0.0, solver='liblinear', max_iter=10000)
log_l2.fit(X_train, y_train)

print("L1 Train Accuracy:", accuracy_score(y_train, log_l1.predict(X_train)))
print("L1 Test Accuracy:", accuracy_score(y_test, log_l1.predict(X_test)))

print("L2 Train Accuracy:", accuracy_score(y_train, log_l2.predict(X_train)))
print("L2 Test Accuracy:", accuracy_score(y_test, log_l2.predict(X_test)))

Best Parameters: {'C': 10, 'solver': 'liblinear'}
Tuned Logistic Regression
-----
```

```

y_train_pred = best_log.predict(X_train)
y_test_pred = best_log.predict(X_test)

print("Tuned Logistic Regression")
print("Train Accuracy:", accuracy_score(y_train, y_train_pred))
print("Test Accuracy:", accuracy_score(y_test, y_test_pred))

# L1 Logistic Regression
log_l1 = LogisticRegression(C=log_cv.best_params_['C'], l1_ratio = 1.0, solver='liblinear', max_iter=10000)
log_l1.fit(X_train, y_train)

# L2 Logistic Regression
log_l2 = LogisticRegression(C=log_cv.best_params_['C'], l1_ratio = 0.0, solver='liblinear', max_iter=10000)
log_l2.fit(X_train, y_train)

print("L1 Train Accuracy:", accuracy_score(y_train, log_l1.predict(X_train)))
print("L1 Test Accuracy:", accuracy_score(y_test, log_l1.predict(X_test)))

print("L2 Train Accuracy:", accuracy_score(y_train, log_l2.predict(X_train)))
print("L2 Test Accuracy:", accuracy_score(y_test, log_l2.predict(X_test)))

Best Parameters: {'C': 10, 'solver': 'liblinear'}
Tuned Logistic Regression
Train Accuracy: 0.9692307692307692
Test Accuracy: 0.956140350877193
L1 Train Accuracy: 0.9824175824175824
L1 Test Accuracy: 0.9736842105263158
L2 Train Accuracy: 0.9692307692307692
L2 Test Accuracy: 0.956140350877193

```

Task 2, Step 3: Regularization Experiments (L1 vs L2)

```

[10]: import matplotlib.pyplot as plt

# Coefficients
coeff_l1 = pd.Series(log_l1.coef_[0], index=load_breast_cancer().feature_names)
coeff_l2 = pd.Series(log_l2.coef_[0], index=load_breast_cancer().feature_names)

print("L1 coefficients:")
print(coeff_l1)

print("\nL2 coefficients:")
print(coeff_l2)

# Count number of non-zero coefficients
print("\nNon-zero L1 coefficients:", np.sum(coeff_l1 != 0))
print("Non-zero L2 coefficients:", np.sum(coeff_l2 != 0))

print("Accuracy Comparison: ")
results = pd.DataFrame({
    'Model': ['L1', 'L2'],
    'Train Accuracy': [
        accuracy_score(y_train, log_l1.predict(X_train)),
        accuracy_score(y_train, log_l2.predict(X_train))
    ],
    'Test Accuracy': [
        accuracy_score(y_test, log_l1.predict(X_test)),
        accuracy_score(y_test, log_l2.predict(X_test))
    ]
})
print(results)

plt.figure(figsize=(12,5))

```

```

        ]
))

print(results)

plt.figure(figsize=(12,5))

plt.plot(log_l1.coef_[0], label="L1")
plt.plot(log_l2.coef_[0], label="L2")

plt.legend()
plt.title("Logistic Regression Coefficients (L1 vs L2)")
plt.show()

L1 coefficients:
mean radius           1.549423
mean texture          0.179249
mean perimeter         -0.036567
mean area              -0.011648
mean smoothness        0.000000
mean compactness       0.000000
mean concavity         0.000000
mean concave points   -14.933701
mean symmetry          0.000000
mean fractal dimension 0.000000
radius error           0.000000
texture error          3.208413
perimeter error        -0.868620
area error              -0.091619
smoothness error       0.000000
compactness error      0.000000
concavity error        2.407315
concave points error   0.000000
symmetry error          0.000000
fractal dimension error 0.000000
worst radius            0.824287
worst texture           -0.599422
worst perimeter         0.133899
worst area               -0.026407
worst smoothness        0.000000
worst compactness       0.000000
worst concavity         -2.334781
worst concave points   -30.911401
worst symmetry           -6.621193
worst fractal dimension 0.000000
dtype: float64

L2 coefficients:
mean radius           4.487493
mean texture          0.271918
mean perimeter         -0.519487
mean area              -0.007428
mean smoothness        -0.721462
mean compactness       -0.694912
mean concavity         -1.741487
mean concave points   -1.643991
mean symmetry          -0.891071
mean fractal dimension 0.036978
radius error           -0.315300
texture error          3.355296
perimeter error        -0.887042
area error              -0.073348
smoothness error       -0.110972
compactness error      0.800577
concavity error         0.922824
concave points error   -0.116166
symmetry error          -0.024919
fractal dimension error 0.146145
worst radius            0.518630
worst texture           -0.642746
worst perimeter         0.164799
worst area               -0.027692

```

```
worst symmetry      -0.0224429  
worst fractal dimension  0.0000000  
dtype: float64
```

```
L2 coefficients:  
mean radius          4.487493  
mean texture          0.271918  
mean perimeter        -0.519487  
mean area             -0.007428  
mean smoothness       -0.721462  
mean compactness      -0.694912  
mean concavity        -1.741487  
mean concave points   -1.643991  
mean symmetry         -0.891071  
mean fractal dimension 0.036978  
radius error          -0.315300  
texture error          3.355296  
perimeter error       -0.887842  
area error             -0.073348  
smoothness error      -0.110972  
compactness error     0.800577  
concavity error       0.922824  
concave points error  -0.116166  
symmetry error        -0.024919  
fractal dimension error 0.146145  
worst radius           0.518530  
worst texture          -0.642746  
worst perimeter        0.164799  
worst area              -0.027692  
worst smoothness       -1.403819  
worst compactness      -1.613764  
worst concavity        -2.977611  
worst concave points   -2.601949  
worst symmetry          -2.994386  
worst fractal dimension -0.026472  
dtype: float64
```

```
Non-zero L1 coefficients: 16  
Non-zero L2 coefficients: 30  
Accuracy Comparison:  
Model Train Accuracy Test Accuracy  
0 L1 0.982418 0.973684  
1 L2 0.969231 0.956148
```

Logistic Regression Coefficients (L1 vs L2)

