

2. Custom Vs Scikit Learn Built Decision Tree

Step 1 : Building a Custom Decision Tree with Information Gain:

```
[6]: class CustomDecisionTree:
    def __init__(self, max_depth=None):
        """
        Initializes the decision tree with the specified maximum depth.
        Parameters:
        max_depth (int, optional): The maximum depth of the tree. If None, the tree is expanded until all
        leaves are pure or contain fewer than the minimum samples required to split.
        """
        self.max_depth = max_depth
        self.tree = None

    def fit(self, X, y):
        """
        Trains the decision tree model using the provided training data.
        Parameters:
        X (array-like): Feature matrix (n_samples, n_features) for training the model.
        y (array-like): Target labels (n_samples,) for training the model.
        """
        self.tree = self._build_tree(X, y)

    def _build_tree(self, X, y, depth=0):
        """
        Recursively builds the decision tree by splitting the data based on the best feature and threshold
        .
        Parameters:
        X (array-like): Feature matrix (n_samples, n_features) for splitting.
        y (array-like): Target labels (n_samples,) for splitting.
        depth (int, optional): Current depth of the tree during recursion.
        Returns:
        dict: A dictionary representing the structure of the tree, containing the best feature index,
        threshold, and recursive tree nodes.
        """
        num_samples, num_features = X.shape
        unique_classes = np.unique(y)

        # Stopping conditions: pure node or reached max depth
        if len(unique_classes) == 1:
            return {'class': unique_classes[0]}
        if num_samples == 0 or (self.max_depth and depth >= self.max_depth):
            return {'class': np.bincount(y).argmax()}

        # Find the best split based on Information Gain (using Entropy)
        best_info_gain = -float('inf')
        best_split = None
        for feature_idx in range(num_features):
            thresholds = np.unique(X[:, feature_idx])
            for threshold in thresholds:
                left_mask = X[:, feature_idx] <= threshold
                right_mask = ~left_mask
                left_y = y[left_mask]
                right_y = y[right_mask]

                info_gain = self._information_gain(y, left_y, right_y)

                if info_gain > best_info_gain:
                    best_info_gain = info_gain
                    best_split = {
                        'feature_idx': feature_idx,
                        'threshold': threshold,
```

```

        if info_gain > best_info_gain:
            best_info_gain = info_gain
            best_split = (
                'feature_idx': feature_idx,
                'threshold': threshold,
                'left_y': left_y,
                'right_y': right_y,
            )

    if best_split is None:
        return {'class': np.bincount(y).argmax()}

    # Recursively build the Left and right subtrees
    left_tree = self._build_tree(X[best_split['left_y']], best_split['left_y'], depth + 1)
    right_tree = self._build_tree(X[best_split['right_y']], best_split['right_y'], depth + 1)

    return {'feature_idx': best_split['feature_idx'], 'threshold': best_split['threshold'], 'left_tree': left_tree, 'right_tree': right_tree}

def _information_gain(self, parent, left, right):
    """
    Computes the Information Gain between the parent node and the left/right child nodes.
    Parameters:
    parent (array-like): The labels of the parent node.
    left (array-like): The labels of the left child node.
    right (array-like): The labels of the right child node.
    Returns:
    float: The Information Gain of the split.
    """
    parent_entropy = self._entropy(parent)
    left_entropy = self._entropy(left)
    right_entropy = self._entropy(right)

    # Information Gain = Entropy(parent) - (weighted average of Left and right entropies)
    weighted_avg_entropy = (len(left) / len(parent)) * left_entropy + (len(right) / len(parent)) * right_entropy
    return parent_entropy - weighted_avg_entropy

def _entropy(self, y):
    """
    Computes the entropy of a set of labels.
    Parameters:
    y (array-like): The labels for which entropy is calculated.
    Returns:
    float: The entropy of the labels.
    """
    # Calculate the probability of each class
    class_probs = np.bincount(y) / len(y)

    # Compute the entropy using the formula: -sum(p * log2(p))
    return -np.sum(class_probs * np.log2(class_probs + 1e-9)) # Added small epsilon to avoid log(0)

def predict(self, X):
    """
    Predicts the target labels for the given test data based on the trained decision tree.
    Parameters:
    X (array-like): Feature matrix (n_samples, n_features) for prediction.
    Returns:
    list: A list of predicted target labels (n_samples,).
    """
    return [self._predict_single(x, self.tree) for x in X]

def _predict_single(self, x, tree):
    """
    Recursively predicts the target label for a single sample by traversing the tree.
    Parameters:
    """

```

```

    return [self._predict_single(x, self.tree) for x in X]

def _predict_single(self, x, tree):
    """
    Recursively predicts the target label for a single sample by traversing the tree.
    Parameters:
    x (array-like): A single feature vector for prediction.
    tree (dict): The current subtree or node to evaluate.
    Returns:
    int: The predicted class label for the sample.
    """
    if 'class' in tree:
        return tree['class']

    feature_val = x[tree['feature_idx']]
    if feature_val <= tree['threshold']:
        return self._predict_single(x, tree['left_tree'])
    else:
        return self._predict_single(x, tree['right_tree'])

```

▼ Step 2 : Load and Split the Iris Datasets:

```

[7]: # Necessary Imports
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score

# Load the Iris dataset
data = load_iris()
X = data.data
y = data.target

# Split into training and test sets (80% training, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

```

Step 3 : Train and Evaluate a Custom Decision Tree:

```

[8]: # Train the custom decision tree
custom_tree = CustomDecisionTree(max_depth=3)
custom_tree.fit(X_train, y_train)

# Predict on the test set
y_pred_custom = custom_tree.predict(X_test)

# Calculate accuracy
accuracy_custom = accuracy_score(y_test, y_pred_custom)
print(f"Custom Decision Tree Accuracy: {accuracy_custom:.4f}")

Custom Decision Tree Accuracy: 0.8000

```

Step 4 : Train and Evaluate a Scikit Learn Decision Tree:

```

[9]: # Train the Scikit-Learn decision tree
sklearn_tree = DecisionTreeClassifier(max_depth=3, random_state=42)
sklearn_tree.fit(X_train, y_train)

# Predict on the test set
y_pred_sklearn = sklearn_tree.predict(X_test)

# Calculate accuracy
accuracy_sklearn = accuracy_score(y_test, y_pred_sklearn)
print(f"Scikit-learn Decision Tree Accuracy: {accuracy_sklearn:.4f}")

Scikit-learn Decision Tree Accuracy: 1.0000

```

Step 5 : Result Comparision:

```
[10]: print(f"Accuracy Comparison:")
print(f"Custom Decision Tree: {accuracy_custom:.4f}")
print(f"Scikit-learn Decision Tree: {accuracy_sklearn:.4f}")

Accuracy Comparison:
Custom Decision Tree: 0.8000
Scikit-learn Decision Tree: 1.0000
```

3. Exercise - Ensemble Methods and Hyperparameter Tuning:

▼ 1. Implement Classification Models:

Step 1: Loading and splitting the data

```
[23]: from sklearn.datasets import load_wine
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import f1_score

# Load dataset
X, y = load_wine(return_X_y=True)

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)
```

Step 2: Training Decision Tree Classifier:

```
[24]: dt_clf = DecisionTreeClassifier(random_state=42)
dt_clf.fit(X_train, y_train)

y_pred_dt = dt_clf.predict(X_test)
f1_dt = f1_score(y_test, y_pred_dt, average='weighted')
```

Step 3: Training Random Forest Classifier:

```
[25]: rf_clf = RandomForestClassifier(random_state=42)
rf_clf.fit(X_train, y_train)

y_pred_rf = rf_clf.predict(X_test)
f1_rf = f1_score(y_test, y_pred_rf, average='weighted')
```

Step 4: Comparing F1 Scores:

```
[26]: print("Decision Tree F1 Score:", f1_dt)
print("Random Forest F1 Score:", f1_rf)

Decision Tree F1 Score: 0.9449614374099499
Random Forest F1 Score: 1.0
```

2. Hyperparameter Tuning:

Step 1: Identifying the three hyperparameters of Random Tree Classifier:

1. `n_estimators` - number of trees
2. `max_depth` - maximum tree depth
3. `min_sample_split` - minimum samples to split a node

Step 2: Using GridSearchCV to optimize these hyperparameters:

```
[27]: from sklearn.model_selection import GridSearchCV

param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 5, 10],
    'min_samples_split': [2, 5, 10]
}

grid_search = GridSearchCV(
    estimator=RandomForestClassifier(random_state=42),
    param_grid=param_grid,
    scoring='f1_weighted',
    cv=5,
    n_jobs=-1
)

grid_search.fit(X_train, y_train)

best_rf_clf = grid_search.best_estimator_

print("Best Parameters:", grid_search.best_params_)

y_pred_best = best_rf_clf.predict(X_test)
print("Optimized F1 Score:", f1_score(y_test, y_pred_best, average='weighted'))
```

Best Parameters: {'max_depth': None, 'min_samples_split': 2, 'n_estimators': 50}
Optimized F1 Score: 1.0

3. Implement Regression Model:

Step 1: Preparing Regression Data:

```
[28]: from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error

# Use alcohol feature as target
X_reg = X[:, 1:]
y_reg = X[:, 0]

Xr_train, Xr_test, yr_train, yr_test = train_test_split(X_reg, y_reg, test_size = 0.2, random_state = 42)
```

Step 2: Training Decision Tree Regressor:

```
[29]: dt_reg = DecisionTreeRegressor(random_state=42)
dt_reg.fit(Xr_train, yr_train)

pred_dt_reg = dt_reg.predict(Xr_test)
mse_dt = mean_squared_error(yr_test, pred_dt_reg)
```

Step 3: Training Random Forest Regressor:

```
[30]: rf_reg = RandomForestRegressor(random_state=42)
rf_reg.fit(Xr_train, yr_train)

pred_rf_reg = rf_reg.predict(Xr_test)
mse_rf = mean_squared_error(yr_test, pred_rf_reg)

print("Decision Tree MSE:", mse_dt)
print("Random Forest MSE:", mse_rf)
```

Decision Tree MSE: 0.3119722222222226
Random Forest MSE: 0.1542667299999946

Step 4: Hyperparameter Tuning using RandomSearchCV

```
[31]: from sklearn.model_selection import RandomizedSearchCV

param_dist = {
    'n_estimators': [50, 100, 200, 300],
    'max_depth': [None, 5, 10, 20],
    'min_samples_leaf': [1, 2, 4]
}

random_search = RandomizedSearchCV(
    estimator=RandomForestRegressor(random_state=42),
    param_distributions=param_dist,
    n_iter=10,
    scoring='neg_mean_squared_error',
    cv=5,
    random_state=42,
    n_jobs=-1
)

random_search.fit(Xr_train, yr_train)

best_rf_reg = random_search.best_estimator_

pred_best_rf = best_rf_reg.predict(Xr_test)
mse_best = mean_squared_error(yr_test, pred_best_rf)

print("Best Parameters:", random_search.best_params_)
print("Optimized Random Forest MSE:", mse_best)

Best Parameters: {'n_estimators': 200, 'min_samples_leaf': 1, 'max_depth': 10}
Optimized Random Forest MSE: 0.151070652463534
```